



Master in High Performance Computing



Istituto Officina
dei Materiali



MPI-IO

- Stefano Cozzini
- CNR-IOM and eXact lab srl

MPI-I/O

- I/O interface specification for use in MPI apps
- Available in MPI-2.0 standard on
- Data model is a stream of bytes in a file
 - Same as POSIX and stdio
- Features:
 - Noncontiguous I/O with MPI datatypes and file views
 - Collective I/O
 - Nonblocking I/O
 - Fortran/C bindings (and additional languages)
- API has a large number of routines..

NOTE: you simply compile and link as you would any normal MPI program.

Why MPI is good for I/O ?

- Writing is like sending a message and reading is like receiving one.
- Any parallel I/O system will need to
 - define collective operations (*MPI communicators*)
 - define noncontiguous data layout in memory and file (*MPI datatypes*)
 - Test completion of nonblocking operations (*MPI request objects*)
- i.e., lots of MPI-like machinery needed

Parallel I/O in MPI

- Why do I/O in MPI?
 - Why not just POSIX?
 - Parallel performance
 - Single file (instead of one file / process)
- MPI has replacement functions for POSIX I/O
 - Provides migration path
- Multiple styles of I/O can all be expressed in MPI
 - Including some that cannot be expressed without MPI

The basic: an example

Just like POSIX I/O, you need to

- Open the file
- Read or Write data to the file
- Close the file

In MPI, these steps are almost the same:

- Open the file: `MPI_File_open`
- Write to the file: `MPI_File_write`
- Close the file: `MPI_File_close`

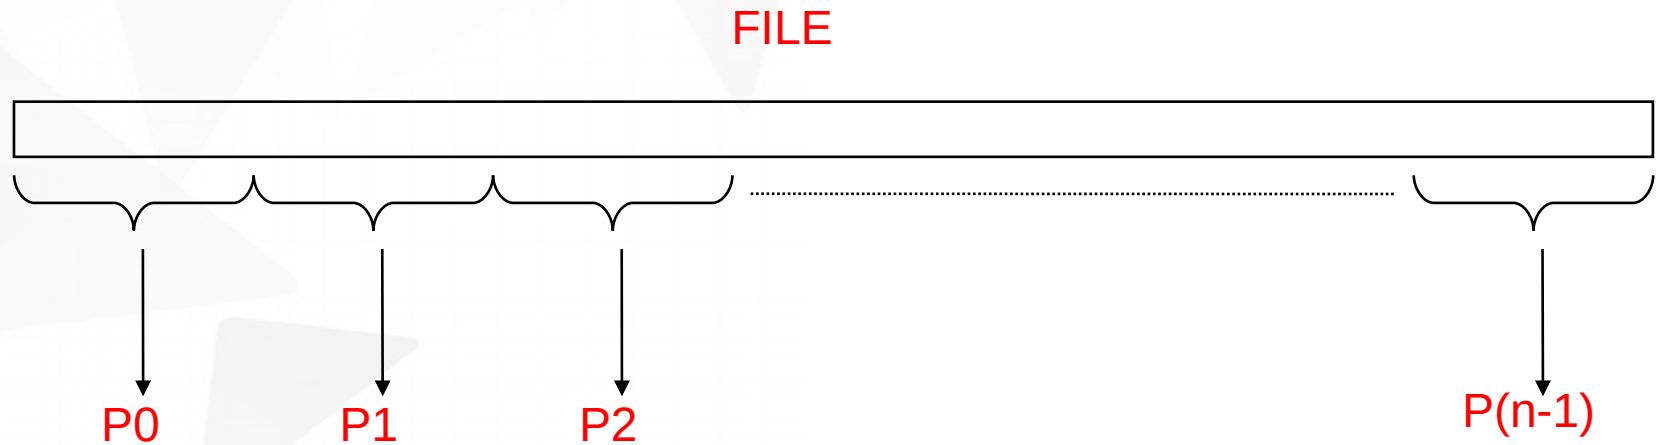
A simple C example

```
#include <stdio.h>
#include "mpi.h"
int main(int argc, char *argv[])
{
    MPI_File fh;
    int buf[1000], rank;
    MPI_Init(0,0);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_File_open(MPI_COMM_WORLD, "test.out", MPI_MODE_CREATE | MPI_MODE_WRONLY,
    MPI_INFO_NULL, &fh);
    if (rank == 0) {MPI_File_write(fh, buf, 1000, MPI_INT, MPI_STATUS_IGNORE);}
    MPI_File_close(&fh);
    MPI_Finalize();
    return 0;
}
```

Some comments

- File Open is collective over the communicator
 - Will be used to support collective I/O, which is important for performance
 - Modes similar to Unix open
 - MPI_Info provides additional hints for performance
- File Write is independent (hence the test on rank)
 - Many important variations covered in later slides
- File close is collective; similar in style to `MPI_Comm_free`

What MPI-I/O is dealing with...



Each process needs to read
a chunk of data from a common file

How to do that ?

```
/* from Gropp's book: page 189*/
#include "mpi.h"
#define FILESIZE (1024*1024)
Int main(int argc, char **argv)
{
    int *buf,rank,nprocs, nints, bufsize;
    MPI_File fh;
    MPI_Status status;

    MPI_Init(&argc,&argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

    bufsize = FILESIZE/nprocs;
    buf= (int *) malloc(bufsize);
    nints = bufsize/sizeof(int);

    MPI_File_open(MPI_COMM_WORLD, "/pfs/datafile",MPI_MODE_RDONLY,MPI_INFO_NULL,&fh);
    MPI_File_seek(fh, rank * bufsize, MPI_SEEK_SET);
    MPI_File_read(fh, buf, nints, MPI_INT, &status);
    MPI_File_close(&fh);

    free(buf);
    MPI_Finalize();
    return 0;
}
```

MPI-I/O API Opening/Closing file

```
int MPI_File_open(MPI_Comm comm,  
char *filename, int amode, MPI_Info info,  
MPI_File *fh)
```

- Collective operation
 - All processes have to provide the same amode
 - comm must be an intra-communicator
- To close the file:

```
int MPI_File_close(MPI_File *fh)
```

MPI-I/O amode values

- `MPI_MODE_RDONLY`: read only
 - `MPI_MODE_WRONLY`: write only
 - `MPI_MODE_RDWR`: read and write
 - `MPI_MODE_CREATE`: create file if it doesn't exist
 - `MPI_MODE_EXCL`: error if creating file that already exists
 - `MPI_MODE_DELETE_ON_CLOSE`: delete file on close
 - `MPI_MODE_UNIQUE_OPEN`: file will not be concurrently opened elsewhere
 - `MPI_MODE_SEQUENTIAL`: file will only be accessed sequentially
 - `MPI_MODE_APPEND`: set initial position of all file pointers to end of file
-
- Combination of several amodes possible, e.g
 - – C: `(MPI_MODE_CREATE | MPI_MODE_WRONLY)`
 - – Fortran: `MPI_MODE_CREATE + MPI_MODE_WRONLY`

Some more observations

- Collective operations across processes within an MPI communicator.
 - Filename must be unique for all processes.
 - Process-local files can be opened with `MPI_COMM_SELF`.
- Initially, all processes view the file as a **linear byte stream**, and each process views data in its own native representation.
 - The file view can be changed via the `MPI_FILE_SET_VIEW` routine.
- Additional information can be passed to MPI environment via the `MPI_Info` handle.
 - The info argument is used to provide extra information on the file access patterns
 - The constant `MPI_INFO_NULL` can be specified as a value for this argument.

File pointer and offset

- In simple MPI-I/O, each MPI process reads or writes a single block.
- We have three means of positioning where the read or write takes place for each process:
 - **Use individual file pointers:**
 - call MPI_File_seek/read
 - **Calculate byte offsets:**
 - call MPI_File_read_at/File_write_at
 - Access a shared file pointer:
 - call MPI_File_seek_shared/read_shared
- Techniques 1 and 2 are naturally associated with C and Fortran, respectively

MPI-I/O API for reading files

```
int MPI_File_seek(MPI_File fh, MPI_Offset offset,  
int whence)
```

```
int MPI_File_read(MPI_File fh, void *buf,  
int count, MPI_Datatype datatype, MPI_Status *status)
```

whence in **MPI_File_seek** updates the individual file pointer according to

MPI_SEEK_SET: the pointer is set to offset

MPI_SEEK_CUR: the pointer is set to the current pointer position plus offset

MPI_SEEK_END: the pointer is set to the end of file plus offset (offset could negative)

Reading a file by using individual pointers (C code)

```
MPI_File fh;
MPI_Status status;

MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &nprocs);

bufsize = FILESIZE/nprocs;
nints = bufsize/sizeof(int);

MPI_File_open(MPI_COMM_WORLD, "your_filename",
              MPI_MODE_RDONLY, MPI_INFO_NULL, &fh);
MPI_File_seek(fh, rank * bufsize, MPI_SEEK_SET);
MPI_File_read(fh, buf, nints, MPI_INT, &status);
MPI_File_close(&fh);
```

See in your github account for this complete example

Reading a file using explicit offset

```
int MPI_File_read_at(MPI_File fh, MPI_Offset offset,  
void *buf, int count, MPI_Datatype datatype,  
MPI_Status *status)
```


Reading a file using explicit offset (F90)

```
include 'mpif.h'

integer status(MPI_STATUS_SIZE)
integer (kind=MPI_OFFSET_KIND) offset

call MPI_FILE_OPEN(MPI_COMM_WORLD, 'my_output_file', &
                   MPI_MODE_RDONLY, MPI_INFO_NULL, fh, ierr)
nints = FILESIZE / (nprocs*INTSIZE)
offset = rank * nints * INTSIZE
call MPI_FILE_READ_AT(fh, offset, buf, nints,
                     MPI_INTEGER, status, ierr)
call MPI_GET_COUNT(status, MPI_INTEGER, count, ierr)
print *, 'process ', rank, 'read ', count, 'integers'

call MPI_FILE_CLOSE(fh, ierr)
```

See in your github account for this complete example

MPI-I/O API for writing files

```
int MPI_File_write(MPI_File fh, void *buf,  
int count, MPI_Datatype datatype, MPI_Status *status)
```

```
int MPI_File_write_at(MPI_File fh, MPI_Offset offset,  
void *buf, int count, MPI_Datatype datatype,  
MPI_Status *status)
```

- To write a file:
 - Set the appropriate flag/s in MPI_File_open:
MPI_MODE_WRONLY Or MPI_MODE_RDWR and if needed,
MPI_MODE_CREATE
 - Use MPI_File_write or MPI_File_write_at

Write files with offset..

```
#include<stdio.h>
#include "mpi.h"
int main(int argc, char **argv){
int i, rank, size, offset, nints, N=16 ;
MPI_File fhw;
MPI_Status status;
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);
int buf[N];
for ( i=0;i<N;i++){
    buf[i] = i ;
}
offset = rank*(N/size)*sizeof(int);
MPI_File_open(MPI_COMM_WORLD, "datafile",
    & MPI_MODE_CREATE|MPI_MODE_WRONLY, MPI_INFO_NULL, &fhw);
printf("Rank: %d, Offset: %d\n", rank, offset);
MPI_File_write_at(fhw, offset, buf, (N/size), MPI_INT, &status);
MPI_File_close(&fhw);
MPI_Finalize();
return 0;
}
```



See in your github account for this complete example

Summarizing so far:

MPI_File_open

MPI_File_seek

MPI_File_read

MPI_File_write

Individual pointer
functions

MPI_File_close

+

ALL INDEPENDENT I/O OPERATION

MPI_File_read_at

MPI_File_write_at

Explicit offset functions

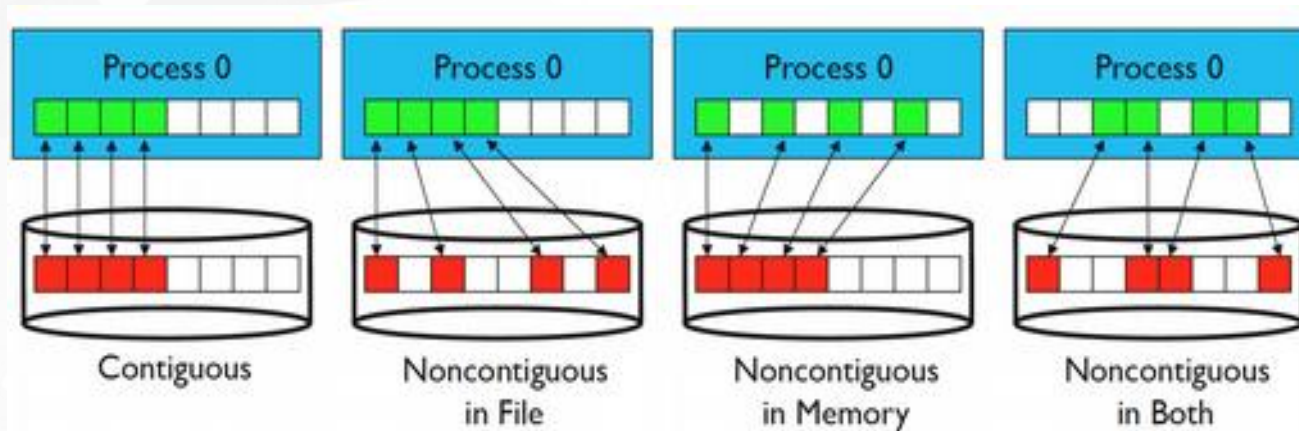
Why do we use independent I/O ?

- Sometimes the synchronization of collective calls is not natural
- Sometimes the overhead of collective calls outweighs their benefits
 - Example: very small I/O duringv header reads

Are we done ?

- YES:
 - These five routines are enough to write any parallel I/O program
- NO:
 - Other MPI-IO routines are for
 - Performance/ Portability/ Convenience
 - Real benefits comes from:
 - non contiguous access
 - collective I/O

Data pattern: contiguous vs non/contiguous



- Contiguous I/O: from a single memory block into a single file region
- Noncontiguous I/O has three forms:
 - Noncontiguous in memory, noncontiguous in file, or noncontiguous in both
- **Structured data** leads naturally to noncontiguous I/O (e.g. block decomposition)
- Describing noncontiguous accesses with a single operation passes more knowledge to I/O system

Considerations on contiguous vs non contiguous

- Best performance comes from situations when the data is accessed contiguously in memory and on disk.
- Commonly, data access is contiguous in memory but noncontiguous on disk.
 - i.e: reconstruct a global data structure via parallel I/O.
- Sometimes, data access may be contiguous on disk but noncontiguous in Memory
 - i.e: writing out a list of neighbors in MD codes.
- A large impact on I/O performance would be observed if data access was noncontiguous both in memory and on disk.

MPI notion of file view

- File view in MPI defines which portion of a file is *visible* to a process
- When a file is first open it is entirely visible to all processes
- The file view of each process can be changed by means of `MPI_File_set_view`
- It is possible to do this operation as many time as you want in a program

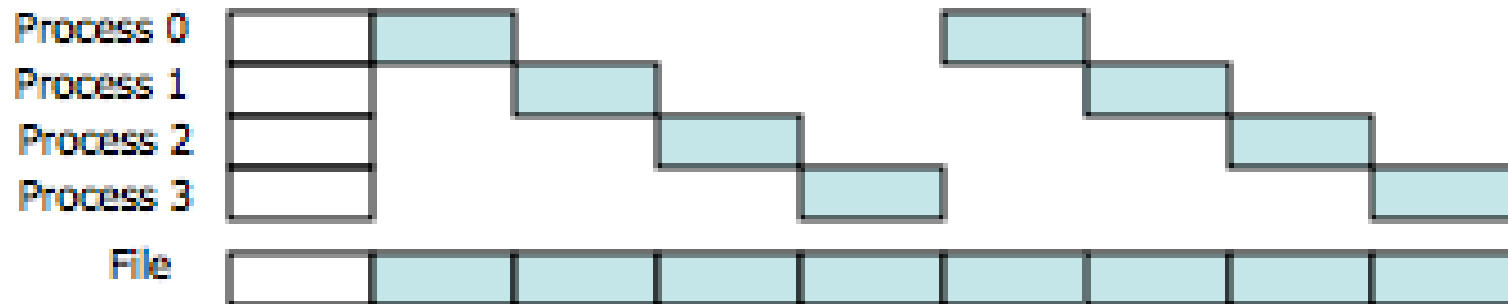

Why do we want to change File view ?

- To indicate which kind of data we are going to read.
 - By default just a bit of stream..
 - We need to use this to ensure portability
- To indicate which part of the file should be skipped... to specify non contiguous access

Using File Views

- Processes write to shared file

FILE



File Views


```
int MPI_File_set_view(MPI_File fh,  
MPI_Offset displacement,  
MPI_Datatype etype, MPI_Datatype filetype,  
char *datarep, MPI_Info info)
```

Specified by a triplet (*displacement*, *etype*, and *filetype*) passed to **MPI_File_set_view**

- *displacement* = number of bytes to be skipped from the start of the file
- *etype* = basic unit of data access (can be any basic or derived datatype)
- *filetype* = specifies which portion of the file is visible to the process (same as etype or derived type consisting of etype)
- Default view: displacement 0 / etype filetype =MPI_BYTE/

Note !

The pattern described by a filetype is repeated, beginning at the displacement, to define the view within the file..



File View

- File view: portion of a file visible to a process
 - Processes can share a common view
 - Views can overlap or be disjoint
 - Views can be changed during runtime
- A process can have multiple instances of a file open using different file views

File View basic example (and usage)

```
MPI_File thefile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
MPI_File_open(MPI_COMM_WORLD, "testfile",
               MPI_MODE_CREATE | MPI_MODE_WRONLY,
               MPI_INFO_NULL, &thefile);
MPI_File_set_view(thefile, myrank * BUFSIZE * sizeof(int),
                  MPI_INT, MPI_INT, "native",
                  MPI_INFO_NULL);
MPI_File_write(thefile, buf, BUFSIZE, MPI_INT,
               MPI_STATUS_IGNORE);
MPI_File_close(&thefile);
```

See in your github account for this complete example

File Interoperability

- Fifth parameter of `MPI_File_set_view` sets the data representation used:
 - **native**: data is stored in a file exactly as it is in memory
 - **internal**: data representation for heterogeneous environments using the same MPI I/O implementation
 - **external32**: portable data representation across multiple platforms and MPI I/O libraries.
- User can also register her own data representation
 - appropriate conversion functions (`MPI_Register_datarep`) should be provided.

Exercises

- Play with the two examples provided (see directory)
- Write two simple MPI programs to write files using `set_view` function accordingly to examples/exercise 1 and 2 of the previous slides.. (optional # 1)

File view example: exercise 1 optional

- Write contiguous data into a contiguous block using file view
- Use derived data type to define filetype in the file view.

P0	1	2	3	4
P1	11	12	13	14
P2	21	22	23	24
P3	31	32	33	34

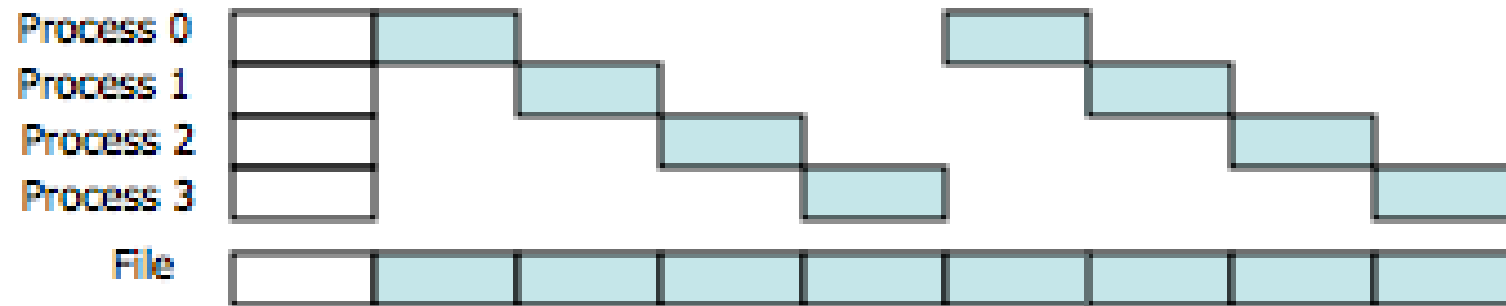
1	2	3	4	11	12	13	14	21	22	23	24	31	32	33	34
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

```
MPI_Type_contiguous(N, MPI_INT, &arraytype);  
MPI_Type_commit(&arraytype);
```

File view Example/exercise 2

- Write a file with the following layout:

1	2	11	12	21	22	31	32	3	4	13	14	23	24	33	34
---	---	----	----	----	----	----	----	---	---	----	----	----	----	----	----



```
MPI_Type_vector(2,, MPI_INT, &fileblk);  
MPI_Type_commit(                                     &fileblk);
```

Links/Reference

- MPI –The Complete Reference vol.2, The MPI Extensions
 - (W.Gropp, E.Lusk et al. -1998 MIT Press)
- Using MPI-2: Advanced Features of the Message- Passing Interface
 - (W.Gropp, E.Lusk, R.Thakur-1999 MIT Press)
- Standard MPI-2.x (or the last MPI-3.x) (<http://www.mpi-forum.org/docs>)
- Users Guide for ROMIO (Thakur, Ross, Lusk, Gropp, Latham)
(<http://www.mcs.anl.gov/research/projects/romio/doc/users-guide.pdf>)
- <http://beige.ucs.indiana.edu/l590/node86.html>