



Master in High Performance Computing



Istituto Officina
dei Materiali



MPI-IO part 2


- Stefano Cozzini
- CNR-IOM and eXact lab srl

Agenda

- Short recap: again on File View
- Collective Operations
- MPI_HINTS
- A final exercise (for MHPC students)


Short recap :

What MPI-I/O is dealing with...



Each process needs to read a chunk of data from a common file

Which kind of data pattern are we performing here ?



But do we need MPI for this ?

- Regular Posix I/O functions can do contiguous access...

- `lseek` (C System Call)

- `lseek` is a system call that moves the read/write pointer of a file to the location of the file. It can be set either in absolute or relative mode.


- Be careful however about `lockf`

interact with the file system, system call `lseek`, exactly what we did on mpi: we can do it without mpi
`lseek` works well

beware of lock mechanism
posix is atomic on some operations, locks them.
in any case just contiguous access:
lock mechanism: atomic operation on the file concurrent access on the file, then processes are queued

the location of
the location can be

What about this case ?



- Simple way:
 - 2 separated calls
 - Read first chunk
 - Read the second chunk

Extremely ine

mpi with different kind of approach than contiguous here non contiguous access of file, can be done on C opening and closing files, very inefficient: because I have to open the file twice, for each reading I read only a small chunk, doubling the latency of my i/o system, get down to the disk for a small amount of data.

if there is a complex non-contiguous pattern it takes ages. You have to reconsider the seek operation. seeking a new position is a very expensive operation.

MPI I/O approach

strong point of
open MPI

- Ability to access NON contiguous data with a single function call

MPI notion of file view

- File view in MPI defines which portion of a file is *visible* to a process
- When a file is first open it is entirely visible to all processes
- The file view of each process can be changed by means of `MPI_File_set_view`
- Read/Write function will be “see” only the visible portion of the file
- `MPI_File_set_view` assigns regions of the file to separate processes

you see non-contiguous access to file, reduce latency with just one command

File Views

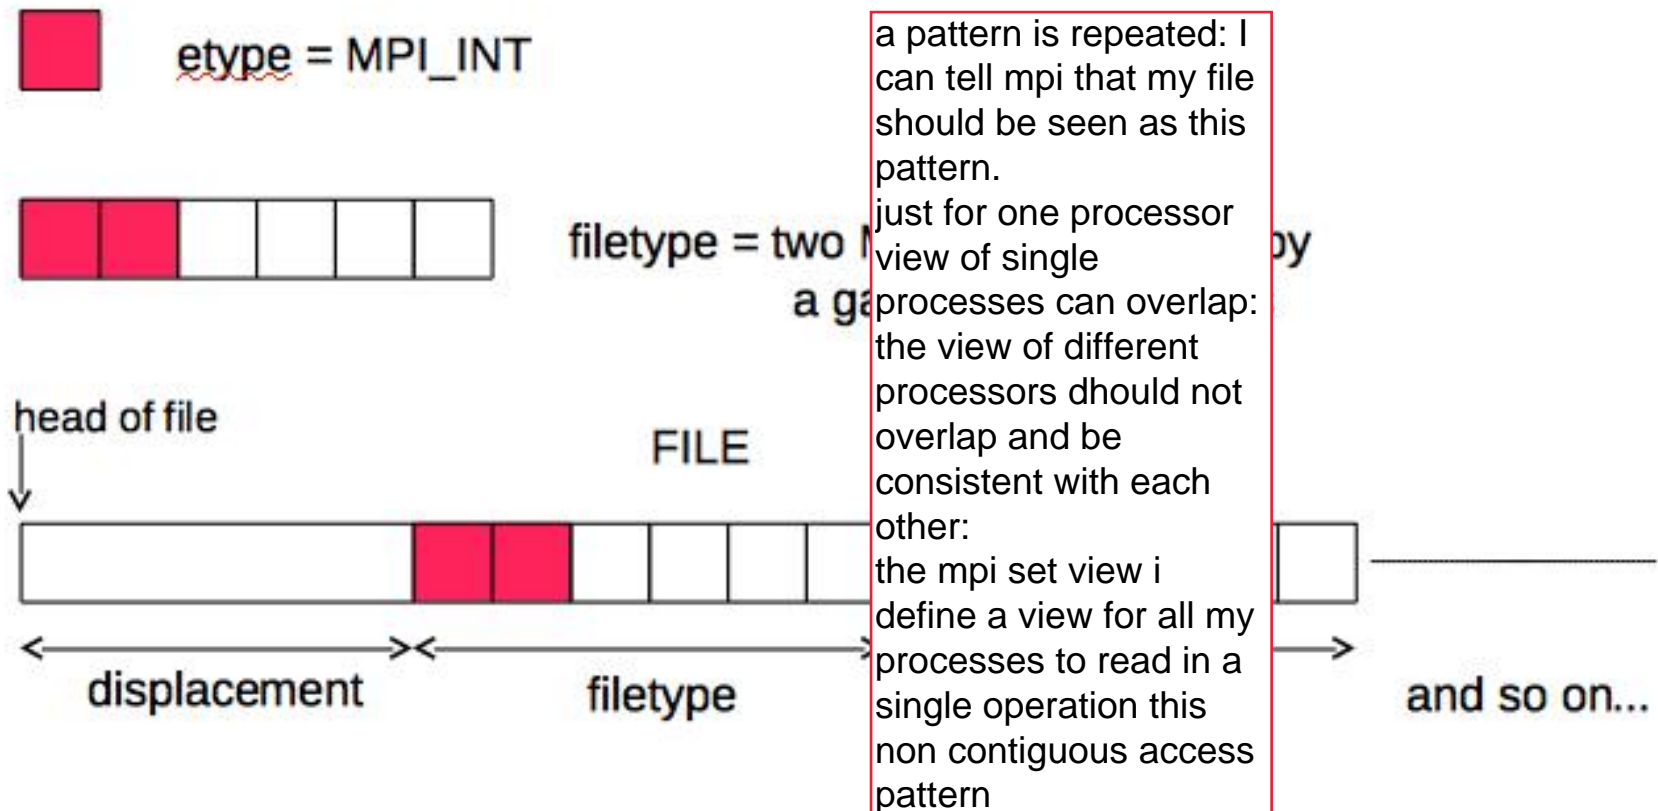
```
int MPI_File_set_view(MPI_File fh,  
MPI_Offset displacement,  
MPI_Datatype etype, MPI_Datatype filetype,  
char *datarep, MPI_Info info)
```

Specified by a triplet (*displacement*, *etype*, and *filetype*) passed to **MPI_File_set_view**

- *displacement* = number of bytes to be skipped from the start of the file
- *etype* = basic unit of data access (can be any basic or derived datatype)
- *filetype* = specifies which portion of the file is visible to the process (same as etype or derived type consisting of etype)
- Default view: displacement 0 / etype filetype =MPI_BYTE/

Note !

The pattern described by a filetype is repeated, beginning at the displacement, to define the view within the file..



File View basic example: contiguous access

```
MPI_File thefile;

for (i=0; i<BUFSIZE; i++)
    buf[i] = myrank * BUFSIZE + i;
MPI_File_open(MPI_COMM_WORLD, "test",
               MPI_MODE_CREATE | MPI_MODE_WRONLY,
               MPI_INFO_NULL, &thefile);
MPI_File_set_view(thefile, myrank * BUFSIZE * sizeof(int),
                  MPI_INT, MPI_INT, "native",
                  MPI_INFO_NULL);
MPI_File_write(thefile, buf, BUFSIZE, MPI_INT,
               MPI_STATUS_IGNORE);
MPI_File_close(&thefile);
```

we open a file:
each processor
sees exactly
what portion of
data they must
see
mpi_file_set_vie
w-->ordered by
rank, mpi int

See in your github account for this complete example

Quick introduction to Derived Data Types

- What are they?
 - Data types built from the basic MPI datatype
Standard defines a general datatype as an
two things:
 - a sequence of basic datatypes
 - a sequence of integer (byte) displacements
- How to use them ?
 - Construct the datatype using a template
 - Allocate the datatype
 - Use the datatype.
 - Deallocate the datatype.

You must construct and allocate a datatype before use
You are not required to use it or deallocate it, but it is

to allow contiguous approach: use
derived data types
(we don't really need so just to show
an alternative to the past slide)
datatypes derived from datatypes
defined by sequences of basic
datatypes, sequence of integer
displacements
each process

castings make the program
unreadable

read integers
optimization not introduced when
developing

don't do sophisticated operations,
don't introduce additional operations,
do them when you need them

premature optimization is the root of
all evils

IO is the slowest operation on the
computer, come to the IO as last
computer turns CPU-bound problem to
IO-bound problem

Main functions

- Datatype constructors:
 - `MPI_TYPE_CONTIGUOUS (count, oldtype, newtype)`
 - Simplest constructor. Makes `count` copies of `oldtype` into `newtype`
 - `MPI_TYPE_VECTOR (count, blocklength, stride, newtype)`
 - Make `count` copies of a block of length `blocklength` with `stride` gaps (stride) in the displacements.
- Allocate and deallocate
 - C
 - `int MPI_Type_commit (MPI_datatype *datatype)`
 - `int MPI_Type_free (MPI_datatype *datatype)`

two different datatypes
strides, regular gaps

`mpi_type_commit`: i can
use the new datatype, free
to deallocate

i define a new datatype from
an old datatype, `newtype`
contiguous types of old
datatype

I take a standard data type
and i make a sequence of
them

Vector, there is stride in
between

stride: space between two
readings

I define a 2d array i wanna
built a datatype to read just
rows, fortran reads by
column,
implementation behind
resolves all location in
memory

`datatype(oldtype)`

`datatype,`

vs for regular

Let us solve exercises

- Take a look at the F90 code where file_set_view routine is introduced
- Compile the code and compare results with output of writeFile_pointer.f90
- Modify the code to use file_set_view using mpi_c_data_type (MPI_TYPE_CONTIGUOUS or MPI_TYPE_VECTOR) and get the same results as the code you start from.

derived datatypes:same thing as building structures in C++
etype specifies the size, the second type is filetype explains the way you are reading
when i use file set view you read an integers and read four of them, specify file type

Exercise 4 optional

- Write contiguous data into a contiguous block using file view
- Use derived data type to define filetype in the file view.

P0	1	2	3	read 4 elements one after the other
P1	11	12	13	14
P2	21	22	23	24
P3	31	32	33	34

1	2	3	4	11	12	13	14	21	22	23	24	31	32	33	34
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

Solution to exercise 4:

contiguous
element is four
elements one
after the other
disp must be
computed in
bytes

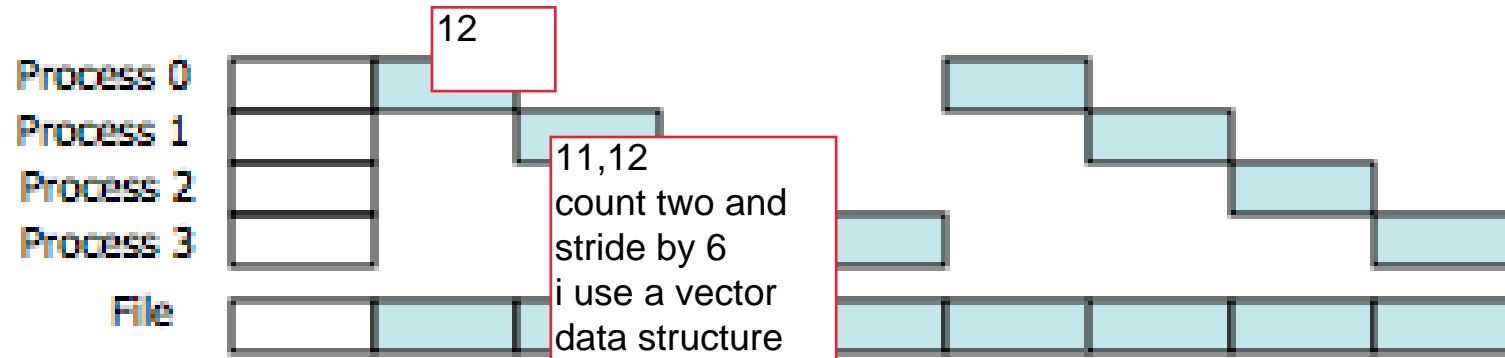
```
integer ierr, i, myrank, BUFSIZE, t  
parameter (BUFSIZE=4)  
  
...  
  
!!define a contiguous derived datatype  
call MPI_TYPE_CONTIGUOUS(BUFSIZE,MPI_INTEGER, filetype,  
ierr)  
call MPI_TYPE_COMMIT(filetype, ierr)  
  
disp = myrank * BUFSIZE * intsize  
  
write(*,*) "myid ",myrank,"disp ", disp  
  
call MPI_FILE_SET_VIEW(thefile, disp, etype, &  
                        filetype, 'native', &  
                        MPI_INFO_NULL, ierr)
```

...

File view example/exercise 4b

- Write a file with the following layout:

1	2	11	12	21	22	31	32	3	4	13	14	23	24	33	34
---	---	----	----	----	----	----	----	---	---	----	----	----	----	----	----



File view example/exercise 4b

```
!!define a different pattern by mea
!!    first parameter: number of gl
!!    second parameter: number of b
!!    third parameter: stride betw
!!
call MPI_TYPE_VECTOR(bufsize/2,npro
    MPI_INTEGER,filetype,ierr)
call MPI_TYPE_COMMIT(filetype, ierr)

disp = (bufsize/2) * myrank * intsi
write(*,*) "myid ",myrank," disp "

...
call MPI_FILE_SET_VIEW(thefile, disp, etype, &
    filetype, 'native', &
    MPI_INFO_NULL, ierr)
```

bufsize=4
i want 2
elements, then i
want two blocks
i have 4
processors and i
divide by 2
now jump away
two elements,
two blocks, now
the stride from
two blocks,
stride is actually
8, between
blocks nprocs,
means we have
4 blocks
between two
blocks

TYPE_VECTOR
nt

procs, &

MPI properties

three main
properties
positioning seek
and whatever
synchronization:
blocking and
non blocking
coordination
collecting
operations
together

POSITIONING

SYNCHRONIZATION

COORDINATION

MP-IO properties : positioning

- Positioning positioning in files
 - Use individual file pointers:
 - call MPI_File_seek/read
 - Calculate byte offsets:
 - call MPI_File_read_at
 - Access a shared file pointer:
 - call MPI_File_seek_shared/read_shared

MP-IO properties : SYNCHRONIZATION

- Synchronization:
 - MPI-2 supports both blocking and non blocking IO routines
 - A blocking IO call will not return until the IO request is completed.
 - A non blocking IO call initiates an IO operation, but not wait for its completion

write in a
blocking and
non blocking
way

mpi_file_iread_a
t is nonblocking

MPI_wait
waiting for
completion of
mpi file write at

Nonblocking I/O

```
MPI_Request request;
MPI_Status status;

MPI_File_iwrite_at(fh, offset, buf, count,
datatype,
                    &request);

for (i=0; i<1000; i++) {
    /* perform computation */
}

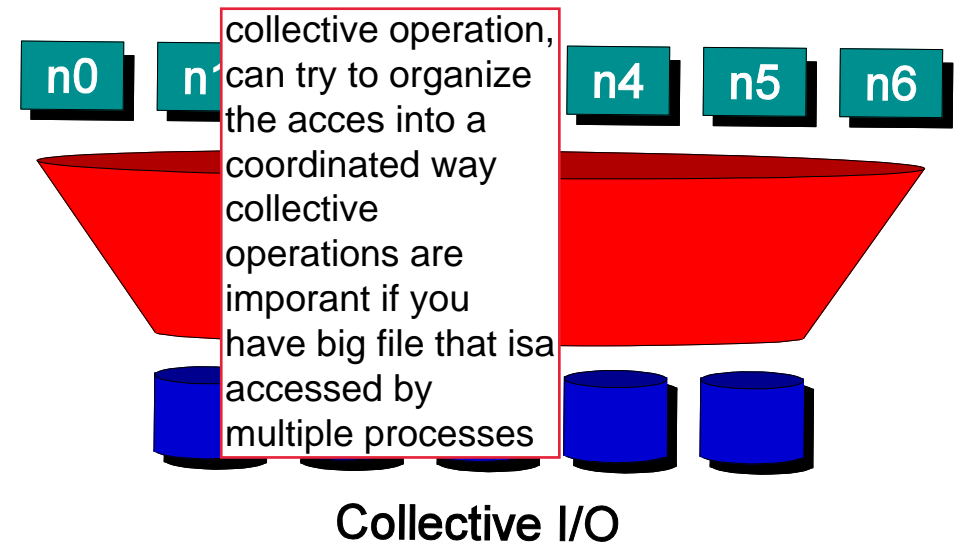

MPI_Wait(&request, &status);
```

MP-IO properties : coordination

when we open a file is a collective operation
every process opens it
coordinates read and write of different
processes. no apparent order to access

- Data access can either take place from individual processes or collectively across a group of processes:
 - collective: MPI coordinates the reads and writes of processes
 - Collective I/O functions must be called by all processes participating in I/O
 - Allows I/O layers to know more about access as a whole
 - independent: no coordination by MPI
 - No apparent order or structure to accesses

Collective I/O operations (1)



Collective I/O operations (2)

pass through `mpi_pi_open` tries to read this file in a coordinated way
attach all at the end:

- **`MPI_File_read_all`**, **`MPI_File_read_at_all`**, etc
- **`_all`** indicates that all processes in the group specified by the communicator passed to **`MPI_File_open`** will call this function
- Each process specifies only its own access information -- the argument list is the same as for the non-collective functions

Collective I/O operations (3)

- Collective I/O is a critical optimization strategy for reading from, and writing to, the parallel system
- The collective read allows all processors to read data together and wait in a coordinated way. This forces all processes in the communicator to read data simultaneously and to wait for each other
- The MPI implementation optimizes the read/write request based on the combined requests of all processes and can merge the requests of different processes for efficiently complete the requests

Optimizing MPI operations..

- Given complete access information, an implementation can perform optimizations such as:
 - **Data Sieving:** Read large chunks and extract what is really Needed
 - **Collective I/O:** Merge requests of different processes into larger requests

1.combine the requests
2.do input output operations in a collective way

read the file in large chunks and throw away what you dont need. I read a lot of stuff and consider a subsection of them

Collective MPI operations: collective buffering

- breaks the IO operation into two stages.
 - first stage uses a subset of MPI tasks (called aggregators) to communicate with the IO servers and read a large chunk of data into a temporary buffer.
 - second stage, the aggregators ship the data from the buffer to its destination among the remaining MPI tasks using point-to-point MPI calls.
- PRO/Cons
 - fewer nodes are communicating with the IO servers, which reduces contention while still attaining high performance through concurrent I/O transfers.
 - Two stages operation: not so easy

a node has a certain number of processors, if i do a collective operation

Collective MPI operations: data sieving

- For independent noncontiguous requests
- ROMIO makes large I/O requests to the file system and, in memory, extracts the data required
- For writing, a read-modify-write is required
- Pro/Cons
 - data is always accessed in large chunks, although at the cost of reading more data than needed. For many common access patterns, the holes between useful data are not unduly large, and the advantage of accessing large chunks far outweighs the cost of reading extra data.
 - In some access patterns, however, the holes could be so large that the cost of reading the extra data outweighs the cost of handling such cases as well.
 - The implementation can decide whether to perform data sieving or access each contiguous data segment separately.

try to read large request from file, then on memory dump to the disk what you need
data accessed in large chunks: read more than what you need

A final summary:

Positioning	Synchronisation	Coordination	
		<i>Noncollective</i>	<i>Collective</i>
<i>Explicit offsets</i>	<i>Blocking</i>	MPI_FILE_READ_AT MPI_FILE_WRITE_AT	MPI_FILE_READ_AT_ALL MPI_FILE_WRITE_AT_ALL
	<i>Non-blocking & split collective</i>	MPI_FILE_IREAD_AT MPI_FILE_IWRITE_AT	MPI_FILE_READ_AT_ALL_BEGIN MPI_FILE_READ_AT_ALL_END MPI_FILE_WRITE_AT_ALL_BEGIN MPI_FILE_WRITE_AT_ALL_END
<i>Individual file pointers</i>	<i>Blocking</i>	MPI_FILE_READ MPI_FILE_WRITE	MPI_FILE_READ_ALL MPI_FILE_WRITE_ALL
	<i>Non-blocking & split collective</i>	MPI_FILE_IREAD MPI_FILE_IWRITE	MPI_FILE_READ_ALL_BEGIN MPI_FILE_READ_ALL_END MPI_FILE_WRITE_ALL_BEGIN MPI_FILE_WRITE_ALL_END
<i>Shared file pointer</i>	<i>Blocking</i>	MPI_FILE_READ_SHARED MPI_FILE_WRITE_SHARED	MPI_FILE_READ_ORDERED MPI_FILE_WRITE_ORDERED
	<i>Non-blocking & split collective</i>	MPI_FILE_IREAD_SHARED MPI_FILE_IWRITE_SHARED	MPI_FILE_READ_ORDERED_BEGIN MPI_FILE_READ_ORDERED_END MPI_FILE_WRITE_ORDERED_BEGIN MPI_FILE_WRITE_ORDERED_END

we didnt discuss this



Right way to access data

Using the Right MPI-IO Function

- Any application as a particular “I/O access pattern” based on its I/O needs
- The same access pattern can be presented to the I/O system in different ways depending on what I/O functions are used and how
- We classify the different ways of expressing I/ O access patterns in MPI-IO into four levels: level 0 – level 3
- We show how the user’s choice of level affects performance

example taken
from the book of
will gropp
chapter 7,
understand how
to read fie

Distribute array access


Large array
distributed
among 16
processes

P0	P1	P2	P3
P4	P5	P6	P7
P8	P9	P10	P11
P12	P13	P14	P15

Each square represents
a subarray in the memory
of a single process

Access Pattern in the file

I have to access
the pattern this
way



Level 0 access

- Each process makes one independent read request for each row in the local array (as in Unix)

I open the file
and each of the
process reads
his portion of
file, everybody
reads by
themselves

```
MPI_File_open(..., file, MPI_MODE_RDONLY, &fh);  
for (i=0; i<n_local_rows; i++)  
    { MPI_File_seek(fh, ...);  
      MPI_File_read(fh, &(A[i][0]), ...); }  
MPI_File_close(&fh);
```

Level 1 access

I can do it in a coordinated way, add an all, if we have some contentions the net effect is sometimes negligible

- Each process make independent read request for each row in the local array (as in Unix) but each process uses collective I/O functions

```
MPI_File_open(..., file, ..., &fh);  
for (i=0; i<n_local_rows; i++)  
    { MPI_File_seek(fh, ...);  
      MPI_File_read_all(fh, &(A[i[0]]), ...); }  
MPI_File_close(&fh);
```

Level 2 access

each process
creates a
derived datatype

no more loop
i can tell read
this file this way
using file set
view

- Each process creates noncontiguous access independent I/O functions

datatype to describe the
defines a file view, and calls

```
MPI_Type_create_subarray(...,  
    &subarray, ...);  
MPI_Type_commit(&subarray);  
MPI_File_open(..., file, ..., &fh);  
MPI_File_set_view(fh, ..., subarray, ...);  
MPI_File_read(fh, A, ...);  
MPI_File_close(&fh);
```

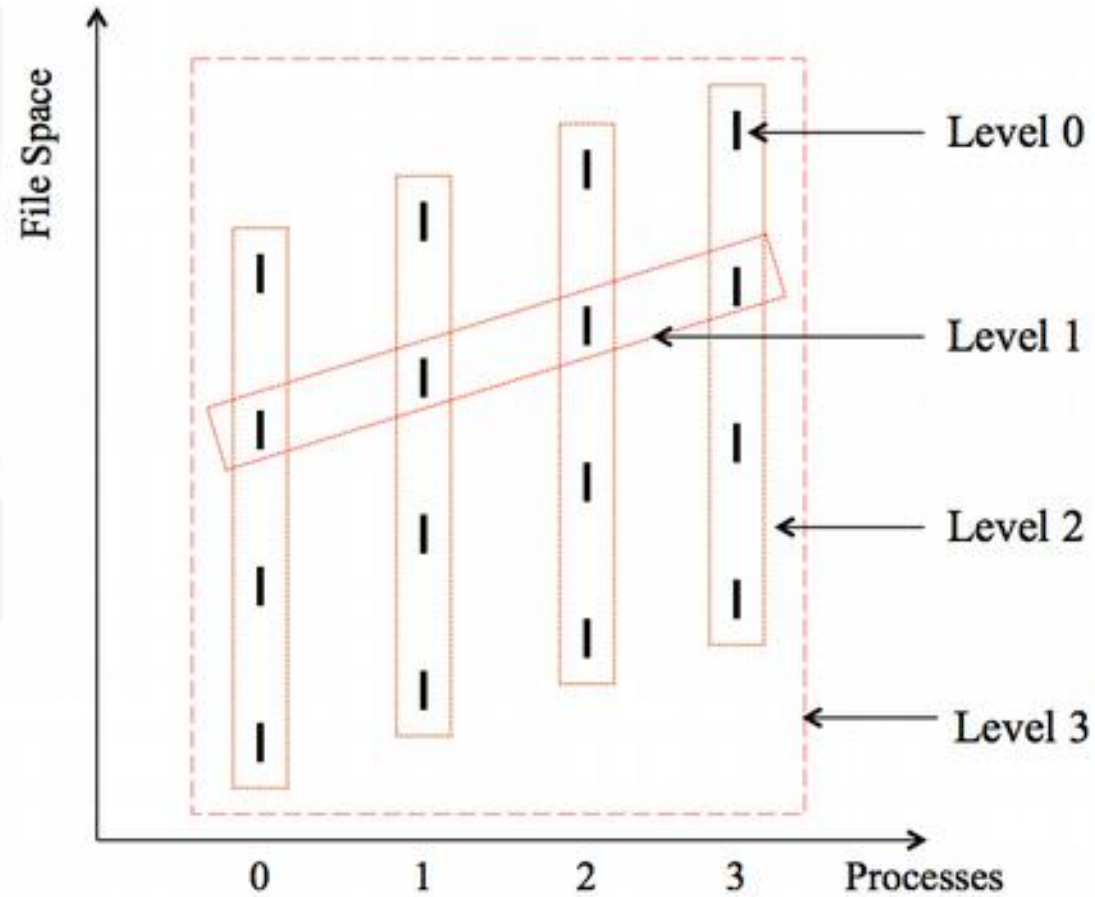
Level 3 access

- Each process creates a derived datatype to describe the noncontiguous access pattern, defines a file view, and calls independent I/O functions

```
MPI_Type_create_subarray(...,  
    &subarray, ...);  
MPI_Type_commit(&subarray);  
MPI_File_open(..., file, ..., &fh);  
MPI_File_set_view(fh, ..., subarray, ...);  
MPI_File_read_all(fh, ...);  
MPI_File_close(&fh);
```

four operations
done altogether

The four level





Conclusions

A very short summary:

- MPI-I/O important features are:
 - The ability to specify noncontiguous accesses
 - The collective I/O functions
 - The ability to pass hints to the implementation

Links/Reference

- MPI –The Complete Reference vol.2, The MPI Extensions (W.Gropp, E.Lusk et al. -1998 MIT Press)
- Using MPI-2: Advanced Features of the Message- Passing Interface (W.Gropp, E.Lusk, R.Thakur-1999 MIT Press)
- Standard MPI-2.x (or the last MPI-3.x) (<http://www.mpi-forum.org/docs>)
- Users Guide for ROMIO (Thakur, Ross, Lusk, Gropp, Latham)
(<http://www.mcs.anl.gov/research/projects/romio/doc/users-guide.pdf>)
- <http://beige.ucs.indiana.edu/l590/node86.html>