
First FHPC Assignment

Nicola Domenis

November 3, 2019

1 PREVIEW

In this assignment we will present the following subjects:

- the analysis of the computational power of our laptop and smartphone;
- the analysis of a strong scaling model for a simple addition problem;
- the resolution of a scalability problem for a simple parallel code that computes pi;
- the parallel implementation of the above addition program;
- the scalability of the addition program.

2 SECTION 0

2.1 Laptop theoretical peak performance

We want to calculate the theoretical peak performance of our own portable computer by using the formula *theoretical peak performance* = clock frequency x FLOPs x number of cores. We gather that *clock frequency* = 2.90Ghz, *FLOPs* = 16 and *number of cores* = 2 for our computer architecture, an intel i7 with a Kaby Lake microarchitecture; thus we compute *theoretical peak performance* = 92.8GFlops/s

	Your model	CPU	Frequency	Number of Cores	Peak Performance
laptop	Asus F556U	Intel Core i7-7500	2.90 GHz	2	92.8 GFLOPs/s

2.2 Smartphone theoretical peak performance

We installed "Mobile Linpack" app and we run a few test. We report here some results, even on repeated trials:

	Model	Sustained performance	Matrix size	Peak performance	Memory
Cellphone	Samsung Galaxy XCover 4	114,81 Mflops/s	250	not calculated	16,00 GB
		145.53 Mflop/s	500		
		157.5 Mflop/s	800		
		201.32 Mflop/s	800		
		155.93 Mflop/s	900		
		109.88 Mflop/s	1000		
		103.14 Mflop/s	2000		

2.3 Laptops,smartphones and the top 500

Let's check now whether our technologies would have competed with the Top500 supercomputers in the past:

	Model	Performance	Top 500 year& position	number 1 HPC system
Smartphone	Samsung Galaxy XCover 4	201,32 Mflop-s/s	does not enter in the top500 of the first year of measurement, the 500th Supercomputer has an Rmax of 0.5 GFlops/s (equal to 2.4 times our smartphone peak performance)	Numerical Wind Tunnel,Fujitsu National Aerospace Laboratory of Japan is first in the year 1993 with a Rmax equal to 124.0 GFlops/s (equal to 616 times our cellphone's sustained peak performance)
Laptop	ASUS F556U	92.8 GFLOP-s/s	3rd position at nov 1993. Remains in the top 10 until nov 1996	We have the same top position with a Rpeak equal to 235.8 GFlops/s(equal to 2.5 times our laptop's theoretical peak performance)

3 SECTION 1

3.1 Model for a serial and parallel summation of n numbers

Here we discuss about modeling a simple program which consists of summing n numbers. A simple pseudocode for the serial program would be:

```
Data:array A[] of values
for i from 1 to n do
    sum = sum + A[i]
end for
return sum
```

If we choose T_{comp} as the time to compute a floating point operation we could calculate the total time of a serial computation as $T_s = N * T_{comp}$,whereas the code simply computes N times(the size of the problem) the sum of two values.

For the parallel program we complicate a little the execution:

```
Data:array A[] of values
Environment: p parallel processors
if Master process then
    Read and Split A[] into p subarrays  $A_i[]$ 
    Send  $p - 1$  subarrays to the other  $p - 1$  processors
    for i from 1 to n/p do
```

```

     $sum_0 = sum_0 + A_0[i]$ 
end for
Collect the resulting  $p - 1$  values  $sum_i$  from the processors
for  $i$  from 1 to  $p$  do
     $sum = sum + sum_i$ 
end for
end if
if Slave process then
    Receive subarrays  $A_i[]$  from the Master process
    for  $i$  from 1 to  $n/p$  do
         $sum_i = sum_i + A_i[i]$ 
    end for
    Send  $sum_i$  back to the Master process
end if
return  $sum$ 

```

If we define the times T_{read} to indicate the time needed to read a variable, and T_{comm} to indicate the time needed to communicate a variable, we can deduce the theoretical execution time of the model:

Read and Split $A[]$ into p subarrays $A_i[]$
 EXECUTION TIME: T_{read}

Send $p - 1$ subarrays to the other $p - 1$ processors
 EXECUTION TIME: $T_{comm} * (p - 1)$

```

for  $i$  from 1 to  $n/p$  do
     $sum_i = sum_i + A_i[i]$ 
end for
    EXECUTION TIME:  $n/p * T_{comp}$ 
    This is a parallel execution, the subarrays are added inside each processor

```

Send sum_i back to the Master process
 EXECUTION TIME: $(p - 1) * T_{comm}$

```

for  $i$  from 1 to  $p$  do
     $sum = sum + sum_i$ 
end for
    EXECUTION TIME:  $(p - 1) * T_{comp}$ 

```

The total sum of the execution times gives $T_p = T_{read} + (p - 1 + n/p) * T_{comp} + 2 * T_{comm}(p - 1)$. We can calculate it with the theoretical values $T_{comp} = 2 \times 10^{-9}$, $T_{read} = 1 \times 10^{-4}$ and $T_{comm} = 1 \times 10^{-6}$

3.2 Scalability of the Model

Once we have the theoretical T_p and T_s we can calculate the Speedup given by the formula $Speedup(p) = T_s / T_p$. We give the following plots on the variable p :




Figure 3.1: Speedup for $N = 10^2$, maximum: $speedup = 0.002$ at $p = 1$




Figure 3.2: Speedup for $N = 10^4$, maximum: $speedup = 0.181$ at $p = 3$




Figure 3.3: Speedup for $N = 10^6$, maximum: $speedup = 8.91$ at $p = 32$




Figure 3.4: Speedup for $N = 10^8$, maximum: $speedup = 147$ at $p = 316$




Figure 3.5: Speedup for $N = 10^{10}$, maximum: $speedup = 905$ at $p = 1000$

We notice that as we increase N we get closer to a case of perfect scaling, on our number of processors. As N grows, adding a certain number of processors accelerates the calculations almost linearly. If N is low, instead, we see that the problem starts scaling around $N = 10^6$, where the speedup grows up to a maximum, and then decreases: the communication time overpowers the advantage of the parallelization, thus lowering the speedup. Simply adding processors to the calculus will not speed up the process because the time it takes to the master node to assign the subarrays to the slaves will grow as well. The algorithm performs well if :

- p corresponds to the maximum speedup;
- the maximum speedup is greater than 1: it makes the parallelization convenient.

4 SECTION 2

4.1 *mpi_pi.c* and *pi.c* execution

We start by executing the two codes *pi.c* and *mpi_pi.c* we have:

```
$ g++ pi.c -o pi.x
$ time ./pi.x 10000000
```

```
# of trials = 10000000 , estimate of pi is 3.141396400
```

```
# walltime : 0.190000000
```

```
real    0m0.275s
user    0m0.271s
sys     0m0.001s
```

And the parallel file:

```
$ mpicc mpi_pi.c -o mpi_pi.x
$ time mpirun -np 10 ./mpi_p.x 10000000
```

```
# walltime on processor 1 : 0.02612305
```

```
# walltime on processor 2 : 0.03022003
```

```
# walltime on processor 3 : 0.02638388
```

```
# walltime on processor 4 : 0.03122497
```

```
# walltime on processor 5 : 0.02647901
```

```
# walltime on processor 6 : 0.02861810
```

```
# walltime on processor 7 : 0.03266811
```

```
# walltime on processor 8 : 0.02701306
```

```
# walltime on processor 9 : 0.03131413
```

of trials = 10000000 , estimate of pi is 3.141720800

walltime on master processor : 0.06575489

```
real    0m1.890s
user    0m11.881s
sys     0m0.630s
```

We should get the longest time of all the parallel execution times of MPI_PLX in order to asses its speed.

Let's collect various run times for a different number of processors.

# of processors	Maximum processor speed
1	0.19799113
2	0.102010919
4	0.05167317
8	0.02782202
16	0.01556611
32	0.04725909
64	0.04147911

We notice that the serial execution time and the single processor parallel execution time differ because of the parallel overhead time: $usertime_{T_p}(1) - usertime_{T_s} + systime_{T_p}(1) = 11.881 - 0.271 + 0.630 = 12.24s$. This is the parallel overhead time for a single process.

Those values are plotted as:




Figure 4.1: Execution time vs number of processors $N = 10^7$

The time decreases inversely to the time. Now lets plot the speedup:




Figure 4.2: Speedup vs number of processors $N = 10^7$

We see that it is not linear: the program scales only until around 20 processors.

Let's repeat our observations by having a larger problem size. Here we have a plot that shows us the maximum execution time against the number of processors. Lets see the case $N = 10^8$




Figure 4.3: Execution time vs number of processors $N = 10^8$

We can see that the speedup decreases for a large number of processors




Figure 4.4: Speedup vs number of processors $N = 10^8$

We can observe that the graph is rising closer to the drawn line, that represents the perfect speedup that is equal to p . Lets see the same graphs for $N = 10^9$:




Figure 4.5: Execution time vs number of processors $N = 10^9$




Figure 4.6: speedup vs number of processors $N = 10^9$

Here the speedup plot is linear, for $0 < p < 64$, thus adding more processors increases linearly the execution time. We have analyzed the strong scalability of the model.

4.2 2.2 Parallel overhead

Here we discuss about identifying a model for deducing the overhead of our program. We study the case where $N = 10^7$ and we plot the difference between the maximum processor walltime and the minimum processor walltime for a different number of processors. We obtain the following graph




Figure 4.7: overhead vs number of processors $N = 10^9$

We can see that the overhead grows as p grows. This is explained by the fact that the overhead increases because of the accessory computation we adopt in order to manage more nodes.

4.3 2.3 Weak scaling

We run a shell script to automatically collect the runtimes for various proportional values of p and N . N is equal to $10^7 * p$ as p grows. We obtain the following plot




Figure 4.8: Weak scaling case: execution time vs number of processors $N = 10^7$

We see that the execution time is almost constant, which is what we expect with a weak scalable program.

5 SECTION 3

Here we have the two c++ codes `SUMNUMBERS_MPI.CC` and `SUMNUMBERS_COLL.CC`.

The two codes implement the pseudocode we wrote above, including the variations on the assignment. Both codes return the sum of N consecutive integers starting from 1 to N and both codes implement the calculation using a parallel approach. Both codes also address the case in

which N is not divisible by the number of processors p . In this case the master node takes care of the summation of the integers that are left out from the slave processors. The code `SUMNUMBERS_MPI.CC` uses only `MPI_SEND()` and `MPI_RECV()` while the code `SUMNUMBERS_COLL.CC` uses the collective operations `MPI_BCAST()` and `MPI_REDUCE()`. The codes were tested using as an input a file containing the number $N = 10^9$.

We collect a few particular times of the execution:

$$\begin{aligned}
 T_{read} &= 3.38554e-05 \text{ seconds} \\
 T_{comp}(N/P) &= 0.355836 \\
 \rightarrow T_{comp} &= 0.355836/100000000 = 3.55836e-9 \\
 T_{comm} &= 2.14577e-06
 \end{aligned} \tag{5.1}$$

We see that these values are close to the theoretical values we gave before, making them plausible.

6 SECTION 4

Now it's time to plot the execution time of the function we wrote, to test the strong scalability. We plot for $N = 10^9$




Figure 6.1: execution time vs number of processors $N = 10^9$




Figure 6.2: speedup vs number of processors $N = 10^9$




Figure 6.3: execution time vs number of processors $N = 10^{10}$




Figure 6.4: speedup vs number of processors $N = 10^{10}$

We see that the problem has gotten worse, it isn't scaling as we predicted in our model. Where by 10^9 the model was perfectly scaling, our model is instead not scaling.

Figure 6.5: speedup vs number of processors $N = 10^{10}$